

Idea

- An edge data center should be able to elect a leader among the other edge data centers in reach (if they meet some arbitrary conditions).
- A leader (that can have a leader as well) can execute tasks on its VM or offload them to its subordinates.
- As last chance the task is offloaded on the Cloud.

Implementation

In my implementation a leader can't be an orchestrator.

This can be easily changed.

At the moment the leader election is done once at the start of the simulation because the conditions on which a device is elected are statics.

This can be changed rescheduling the election with a delay.

LeaderEdgeDevice class

I started by creating a custom class (*LeaderEdgeDevice*) that extends *DefaultDataCenter* using the proposed *ClusterEdgeDevice* as a base.

In **startInternal** a task with tag `LEADER.ELECTION` is scheduled with `INITIALIZATION_TIME + 1` delay.

In process event the custom tag is caught.

There is a check for:

- the device being an edge data center,
- the device being an orchestrator.

This allows data centers not flagged as orchestrators to not have a leader (although possibly being elected as such by other data centers, even that case is easily fixable).

- "LEADER" is the orchestration method,

If everything is true then the method **leader** is called.

- in a loop every datacenter is taken into account and is checked whether it:
 - is not the same data center as the one evaluating,

- is an edge data center
- the distance between the two data centers is smaller than the range of the edge data centers.
- If it is the case then there is an evaluation about the MIPS (I used that criterion for electing a leader).
 - If the MIPS of the candidate are greater than ones of the evaluator
 - and the max MIPS of the data center seen until that point is lower than the MIPS of the candidate

then the candidate becomes the (potentially temporary) leader, the max MIPS seen is set to the new leader's MIPS.

Upon ending the loop if a leader has been found the current datacenter is added to its subordinates list.

Otherwise the orchestrator flag is removed as well as the data center presence in Orchestrators List.

Leader algorithm

In order to find a suitable VM using this hierarchy I created a simulation algorithm called LEADER.

I extended Orchestrator in order to implement a custom findVM method.

If LEADER is used as simulation algorithm there is a check for the - presence of both Cloud and Edge in the simulation architecture. If it is not the case an error is displayed and the simulation stopped.

Otherwise the architecture is restricted to "Edge" and the method leader is called.

Here this order of seek for a suitable VM is followed:

1. Among the hosts of the orchestrator.
2. Among the hosts of the leader of the orchestrator (there should not be a situation where the orchestrator has no leader and is chosen for the previous step so the presence of a leader should be certain at this phase).
3. Among the hosts of the leader's subordinates (excluding the original orchestrator).

If the previous steps should result in no VM selected a try is done with the algorithm INCREASE_LIFETIME (could have been anything else) on the Cloud.

The condition based on which every level is scanned in search for a VM is completely arbitrary.

Particular cases

Chaining

A case of chaining can occur when the data center A has B as leader but B in turn has C as leader.

Both A and B are orchestrator so they can both receive tasks.

During findVM, if a task has A as orchestrator, the task can only be offloaded to B, its subordinates (excluding A) and the Cloud.

This doesn't allow B to offload the original task to C. In fact B act only as a leader in this scenario.

Recursion

A particular case of the situation above is when $A = C$.

This is not a problem and follow the limitation described above.

Isolation

If an edge data center is isolated in such manner that it doesn't have both leader and subordinates it became useless.

Not having a leader elect it as such but it has no subordinates so it won't receive tasks due to the fact that is removed from the orchestrators list.

Parity

In case of condition parity between data centers every device will act as leader, thus not being an orchestrator neither having subordinates.

This could lead to an empty orchestrators list.

Rejection

A datacenter can't reject to be elected as leader.

This will force it to execute a task of its self proclaimed subordinates.

This behavior is easily fixable by placing a check for the value of a custom setting flag "consent" before electing.

Limitations

My actual implementation take as granted the omniscience of the edge data center in regards of its leader and its subordinates.

This is not realistic because the method findVM could discover a VM not owned by the orchestrator.

To simulate this behaviour a custom event could be scheduled to the leader/subordinates simulating the forward of the request.

Due to PES managing orchestration in a stateless way this is possible but only at the cost of invalidating the results of the simulation.

In fact findVM would result in a failure at the eyes of SimulationManager but, in reality, a new schedule for the task is yet to be done following the steps: leader, subordinates and cloud.